

Interview story: FDA behavior-clone system ("parallel")

Use for: "Tell me about a challenging project," "a system you designed end to end,"

"a time you used AI/LLMs in production," "dealing with ambiguity."

The 2-minute version (STAR)

****Situation.**** I wanted a system that tracks changes in the FDA regulatory world and predicts how each stakeholder behaves in response: what enforcement does next, what reviewers will ask for, how industry adapts. News monitoring tells you what happened; nothing told you what happens next.

****Task.**** Build it end to end on parallel.ai's API suite: sense changes, extract structured understanding, build behavior-pattern "clones" of the three stakeholders, simulate their responses side by side, and surface it on a live dashboard.

****Action.****

- Designed a five-stage pipeline: ****Sense**** (Search / Extract / Monitor) ****Clone**** (FindAll + Task for pattern mining, Chat API for the clone runtime) ****Surface**** (FastAPI dashboard, Railway, auto-deploy from GitHub).
- The core design decision: clones are not just personas in a system prompt. I built a ****pattern library**** mining (batch, expensive) vs. matching (deterministic, free) vs. reasoning (LLM) an early version let enforcement patterns leak into the reviewer clone's prompt and the behaviors cross-contaminated. Fixed by partitioning the library per stakeholder in the matcher.
- ****What would you do with more time?**** Live mining loop on a schedule, calibration tracking per pattern (promote/demote confidence by observed accuracy), and a dissent view highlighting where clones disagree.
- ****Ambiguity?**** parallel.ai's docs left several API fields uncertain (monitor payload shape, citation fields, task output shapes). I isolated every uncertainty behind a thin client with defensive parsing and TODO markers,

so verification with
a real key touches one file.