Identity Pack: ML Engineering Manager

Use this pack when interviewing for Engineering Manager / Senior EM roles in machine learning, applied AI, conversational AI, and LLM platform teams.

The one-line identity

"I lead AI engineering teams that ship machine learning systems into production."

Say it in the first sentence of the intro. Every story in the loop should be recognizable as an ML story: a model with behavior, an evaluation loop, production failure modes, iteration.

What this identity emphasizes

- Ownership of models running in production - their behavior, quality, and failure modes
- ML systems work: grounding, guardrails, retrieval quality, evaluation, model-agnostic interfaces
- ML in regulated environments: traceability requirements, model governance, risk frameworks
- Hands-on technical leadership: close enough to the code to debug alongside the team

What this identity avoids

- The three-discipline list ("data science, data engineering, and software engineering"). It splits the panel into "data person" and "full-stack generalist" camps. Never open with it.
- Leading with dashboards, cost savings, or data plumbing. Business outcomes stay, but they follow the ML substance, never lead.
- "I don't go into actual coding." Replace with: "I stay close enough to the code to debug alongside my team" - and prove it in technical rounds.
- Three diffuse reasons for the move. One crisp reason plus the deliberate step-down framing (see level-framing.md).

Files in this pack

- `intro-pitch.md` - the 75-second opener. One identity, first sentence.
- `story-llm-production-app.md` - the LLM application over 40+ manufacturing systems. Grounding, PII guardrails, interface coupling. STAR, ~2 minutes spoken.
- `story-rag-pdf-digitization.md` - the 400K-document RAG system. Chunking decisions, human evaluation, traceability as the product, the risk framework as model governance. STAR, ~2 minutes spoken.
- `level-framing.md` - the deliberate step-down answer, two sentences. Defuses "where do we place him" before the panel has to ask.

How to pick an identity per role

This folder is the first of several. The idea: one folder per role family, each with its own one-line identity, pitch, and reframed stories drawn from the same underlying experience. Future packs (data leadership, full-stack/platform EM) reuse the same stories with different framing - the facts never change, the lens does.